

AutoTriager Shop: Evidence-Grounded Investigation and a Bounded Evaluation

Dahe Chen - CS 5588 Data Science Capstone, Challenge 2

October 1, 2026 - Completed bounded development and confirmation round

1. Abstract

A junior on-call developer new to a shopping system may confuse payment failure with checkout symptoms. AutoTriager provides a read-only investigation feature: select an incident, inspect cited telemetry, and record judgment. Python selects structured evidence; optional Gemini reasoning proposes a service to inspect or defers. Thirteen capture attempts yielded ten accepted windows; nine were eligible for modeling within one development and two confirmation groups. Thirty API calls yielded twenty-eight responses and two HTTP 503 failures. Balanced selection alone established no diagnostic gain. An adaptive investigation-priority change matched the payment label in development, but frozen confirmation selected checkout once and payment once. Both methods deferred on all four confirmation controls. The direct reference matched payment twice, although it asks a different question. Valid citations and control abstention did not ensure faithful explanations. This single-mechanism study establishes an inspectable pipeline, not optimal inspection, causal correctness, generalization, or user benefit.

2. Introduction and research question

Payment failure can produce checkout dependency errors and storefront failure. A junior developer understands application basics but lacks this system's relationships and history. The decision is which service to inspect first, supported by original records.

The semester project concerns incident investigation. Challenge 2 selects P1: load one system's incident, analyze it, inspect sources, and record review. Input is a bounded telemetry interval; output is a candidate or insufficient-evidence decision, references, missing information, and judgment. Historical management and richer interactions follow this first complete path.

The bounded question is: **given payment-failure telemetry with propagated symptoms, can a fixed-budget pipeline propose a checkable first service to inspect while deferring unsupported attributions?** Development first tested initiating-service attribution, then explicitly amended the task to investigation priority. Intervention-label agreement is descriptive: it does not identify the optimal inspection action. The contribution is the inspectable feature and controlled study, with candidate failures and unresolved evidence reported.

3. Related work

[OpenRCA, ICLR 2025](#), Section 2.3, formulates recovery of requested fault attributes from telemetry; RCA-Agent explores records using generated Python. [EviRCA, arXiv:2609.19825v1](#), Sections III-B/C, separates deterministic extraction from bounded read-only reasoning. Its Section V-B discusses assumptions including known fault counts and candidate sets. Thus extraction before reasoning, evidence cards, and read-only analysis are existing ideas. [MicroRCA-Agent, arXiv:2509.15635v1](#), Sections 3.5 and 4.1, provides structured reasoning and reports fabricated call-chain evidence, motivating a separate support audit.

These are closest-work references, not locally replicated baselines. Published benchmark scores are not comparable to this live Shop protocol. The proposed study emphasizes selective investigation priority and incomplete observations; it has not established a literature gap or superiority over these methods.

4. Problem and three-stage design

4.1 Human decisions and honest provenance

Earlier AI-assisted discussions established one system, a small case set, read-only analysis, evidence inspection, and accept/reject/uncertain review. They prioritized existing analysis over agent orchestration. "Human Design" identifies the user's decisions before accepting implementation, not a time-stamped, entirely AI-free original document.

The user's P1 combines incident loading and existing-method analysis, followed by checkable result presentation. This complete path comes first because later features depend on a usable investigation. History, a richer interface, and evaluated analysis improvements are P2; additional systems are P3. The current cycle implements one local path rather than the entire semester roadmap.

Benchmark replay did not establish the chosen shopping scenario. The user retained the investigation boundary but switched to a runnable shopping workflow. Figure 1 reconstructs those recorded decisions.




Figure 1. Reconstructed human-selected workflow from recorded choices. It is not a recovered original drawing.

4.2 AI proposal and critical evaluation

Codex examined the handout, design records, repository, and simulator. The instruction, paraphrased, was: investigate a running shopping incident, isolate labels, cite evidence or abstain, and use a capable direct comparator. This summarizes the actual conversation.

The assistant proposed compact records, stable IDs, timestamps, trace relationships, ranking, optional model interpretation, and a bilingual interface. Separating computation from explanation helped implementation; benchmark framing, weak comparison,

runtime feasibility, and evidence sufficiency still required correction.




Figure 2. Reconstructed AI proposal. This summarizes the earlier direction; it does not imply every proposed component was implemented.

4.3 Reviewed co-design

Human judgment retained the review boundary, accepted structured provenance, rejected benchmark replay as product validation, simplified orchestration, and required private labels and a strong comparator. Development responses demanded deeper-cause evidence although the application intended to prioritize inspection. Co-design explicitly clarified that task while preserving the same evidence. This observation does not establish that the earlier prompt was defective. These choices improve inspectability by design; their user benefit remains untested.

Table 1 is reproduced at full width in Appendix B.

5. Data, method, and application

The official source is [OpenTelemetry Demo 3.1.0](#), commit dedc0178918e260823323b8d95005a8cb924b007, licensed Apache-2.0. Prometheus/Jaeger captures produced three accepted development windows containing 2,793 observations: 393 metrics and 2,400 spans, with no captured logs. The metric `span_error_rate` measures error-span events in calls/s, not a percentage, and is span-derived rather than independent causal corroboration. Table 2 records each window.

The final modeled cohort contains nine dependent phase windows: three development and six confirmation, totaling 8,379 observations (1,179 metrics and 7,200 spans). These are repeated conditions of one mechanism, not nine fault categories. The additional accepted 004 normal is preserved outside the complete-group model cohort.

Observations retain ID, service, UTC time, signal kind, numerical fields, trace/span relationships, and source URLs. Python validates fields, normalizes time, selects records, and ranks explicit anomalies. Public `incident.json` and `observations.json` supply inference. These preserve extracted telemetry fields, not complete original HTTP response archives. Source availability depends on the live system's retention. Private intervention manifests supply evaluation only; known answer fields are rejected from public inputs.




Figure 3. Implemented official capture-to-analysis architecture. Private labels branch only to evaluation; they are excluded from model and UI inputs.

The retrieval unit is an observation, capped at 48 records. Priority sorts explicit `is_anomalous` records first, `ERROR/5xx` records second, and remaining records last, breaking ties by time and ID. The provisional selector round-robins metric, span, and log queues in that order, skipping empty queues and preserving within-kind priority. Official metrics have no `is_anomalous` field, so their order is time/ID rather than measured anomaly magnitude. No embeddings, trained retriever, fine-tuning, or runtime agent team is used. Gemini returns status, candidate, reason, IDs, and missing information. A transparent rule baseline provides another starting point; its heuristic scores are not probabilities.

The application resolves citations and records accept/reject/uncertain review with reasons. Grounded validation requires a candidate present in the selected records, no invalid cited ID, and candidate-local citations from at least two signal kinds; otherwise it defers. The direct validator resolves IDs and checks candidate membership but has no two-kind gate. These structural checks do not establish semantic support. Multiple signal kinds can describe the same request; reference integrity, factual support, and localization remain distinct endpoints.

The user selects a neutral incident alias, reviews provenance, runs the baseline or optional Gemini analysis, opens evidence, and records judgment. API errors display incomplete analysis, not absence of an incident. Language switching preserves case state. Figure 4 shows actual official fault evidence in the application.

Figure 4 appears at full width in Appendix C.

6. Experimental setup

The live deployment completed 18 source image builds and seven third-party pulls, starting 25 containers. Saved image identities matched the prepared images. The source tree remained clean; mutable flags were isolated outside it. Ports bind to localhost: Shop proxy 8080, Prometheus 9090, and Jaeger through proxy 8080. The scheduler was disabled; fixed load settings use five Locust users and HTTP/browser weights 9:1.

The first normal attempt contained checkout/payment ERROR spans and failed the unchanged acceptance gate. Investigation found checkout and catalog Go services limited to 20 MiB, with heavy cgroup pressure. A derived configuration raised each to 128 MiB. At 17:52:23 UTC both were healthy, zero unhealthy containers were recorded, and the recorded full-pressure value was zero. This is a host resource repair, not proof that every earlier error had that cause or a diagnosis improvement.

The pilot uses `paymentFailure=100%` with normal/recovery controls, 180-second warmup and recording windows, 75-second settling, and flag checks every 30 seconds. Acceptance requires checkout-to-payment linkage, visible payment error during fault, and no checkout/payment errors on the specified linked control paths. "Clean" concerns these observed paths, not every service or unobserved request. Gating conditions diagnosis on visible evidence; capture yield must be separate.

One development triplet permits at most three candidate changes. Two later fresh triplets are reserved after freezing method and scorer. Phases, traces, and repeated conditions are dependent; one mechanism tests narrow repeatability, not fault coverage.

Amendment A1, recorded at 18:25 UTC before official calls, addressed modality starvation: all 48 prioritized fault records were spans, conflicting with two-kind validation. Six original-selector and six balanced-selector calls tested this engineering hypothesis. The policy changes composition and ordering together; their effects are not separated. Balanced selection alone established no diagnostic gain.

Amendment A2 at 18:42 UTC followed those responses and simulated review. It allocated the remaining six calls to changing only the grounded objective: choose a defensible inspection priority, without asserting deepest cause. Balanced records/order, model, direct prompt, and validators stayed fixed. The direct prompt was called again and is a descriptive reference, not a deterministic control or same-task superiority comparator. Both metrics-only and spans-only experiments were deferred and remain untested. These are adaptive development amendments, not untouched preregistration.

The recovery timeout prompted explicit `127.0.0.1` origins for replacement `shop-pilot-002-d`. A faster probe does not prove the timeout's cause. The rejected interval is retained; timing, traffic, and gates stayed fixed. Payment failure was restored off and verified before accepted recovery.

Within each condition, prompts share identical records/order: `gemin-3.5-flash-lite`, temperature zero, 2,048 output-token cap, 90-second timeout, seed 5588, alternating order, and no retries. Raw and application decisions are preserved; alternative validators reuse raw responses without calls. Common-validator results separate prompt behavior from postprocessing. Quota/authentication errors stop calls and remain missing analyses.

`selected_method.json` froze balanced selection plus investigation-priority framing at 18:48:30 UTC, including

function/scorer hashes. This is a provisional engineering choice: the full direct-reference no-regression gate is unknown because recovery returned 503. Confirmation triplet 003 started at 18:48:36 and completed three accepted captures at 19:10:24, including verified restoration. Its six calls completed without API failure. Triplet 004 accepted its normal phase but failed during fault collection at 19:22:55 because the feature API timed out; recovery was not attempted. Fault-off restoration was verified. Amendment A3, recorded before replacement 005 started at 19:27:36, permits one entire fresh group with its own baseline and unchanged methods. The remaining six calls are conditional on all replacement phases being accepted. All attempted phases are retained; no further replacement is allowed. This is three attempted confirmation groups, at most two completed and modeled, not an unchanged preregistration. No tuning or held-out selector comparison is allowed.

7. Results and evidence audit

Official event	Recorded outcome	Interpretation
Deployment / preflight	25 containers; 18 source builds, seven pulls	Engineering prerequisite
<code>shop-pilot-001-a</code> normal	Rejected: checkout/payment ERROR spans	Invalid control; no model error
Repair, 17:52:23 UTC	Only checkout/catalog: 20 -> 128 MiB; source clean	Healthy runtime checkpoint
<code>shop-pilot-002-a</code> normal	Accepted 18:02:26; 132 metrics, 800 spans; five linked traces, ten payment spans, zero linked payment ERROR	Window 17:57:09-18:00:21 UTC
<code>shop-pilot-002-b</code> fault	Accepted 18:14:43; 129 metrics, 800 spans; 13 linked traces, 26 payment spans, all 26 ERROR	Window 18:09:26-18:12:38 UTC
<code>shop-pilot-002-c</code> recovery	Rejected 18:15:13: feature-API read timeout	Capture failure retained
<code>shop-pilot-002-d</code> recovery	Accepted; 132 metrics, 800 spans; seven linked traces, 13 payment spans, zero linked payment ERROR	Replacement window 18:28:04-18:31:04 UTC
Confirmation triplet 003	Three accepted phase windows; completed 19:10:24	Six completed calls; findings below
Confirmation triplet 004	Normal accepted; fault timed out; recovery not attempted	Restoration verified; normal retained but unmodeled
Replacement triplet 005	Three accepted windows; completed 19:49:23; restoration verified	Six of six calls complete; no further replacement

Table 2. Final capture accounting. Development yield is 3/5. Confirmation 003 and 005 each add three accepted phases; 004 adds one accepted normal and one failed attempt. Overall yield is 10/13 attempted phases from 14 scheduled phases. 004 recovery was unattempted and is excluded from attempted-capture counts. Its accepted normal is retained but unmodeled under complete-group eligibility. Core-path acceptance is not global absence of errors.

Eighteen scheduled development calls produced sixteen responses and two HTTP 503 failures: original grounded/normal and priority direct/recovery. Neither is an abstention, quota event, or retried success. Table 3 preserves completed denominators and raw/application decisions. There is one development triplet, not eighteen independent incidents.

Table 3 is reproduced at full width in Appendix B.

The sixteen completed calls recorded 115,601 input tokens and 2,593 output tokens, total 118,194. Completed API latency ranged

from 1.090 to 41.895 seconds; the two failed requests lasted 26.673 and 61.466 seconds. These observations do not establish a stable speed ranking or a billed currency cost.

The independent simulated audit in CLAIM_AUDIT_20261001.md separates observable facts from interpretation, using supported, contradicted, insufficient, and not-assessed judgments. It provides locatable qualitative findings, not an expert grounding score.

Case	Expected observable evidence	Selected/cited evidence and audit
Normal 002 - a	Clean observed payment path; other errors may remain	Core rates are zero; load-generator metric-00019 is 0.1666667 calls/s and ERROR spans exist. Priority abstains, but universal zero-error wording is contradicted.
Fault 002 - b	Payment-local error and linked caller evidence	Cites metric-00022 (0.0666667 calls/s, baseline 0), span-00007 and 00008 ERROR. Parent span-00006 is selected but uncited. Observed facts supported; causality unproved.
Recovery 002 - d	Restored core path; preserve residual uncertainty	Core rates are zero; load-generator rate is 0.1166667 calls/s and span-00001 is ERROR without a failing backend link. Priority abstains with empty citations.

Table 4. Expected versus visible evidence. IDs are case-local. Payment/checkout nesting is observed, but the metric derives from spans and does not independently prove cause. The priority fault's missing parent citation limits completeness. Empty-citation controls lack a direct rationale trail. Both incomplete-input views remain untested.

Balanced-only selection showed no diagnostic improvement. Priority's primary development fault/control gates were observed, but the unchanged direct recovery comparison is unavailable. Thus not all retention gates passed. The provisional choice carries this unresolved gate and unfaithful control explanations into confirmation. Failures span capture transport, API availability, selector/validator mismatch, conservative reasoning, and semantic overgeneralization; none should disappear into one accuracy number.

7.1 First frozen confirmation

All six calls on triplet 003 completed. Direct raw and application outputs matched payment on the fault and abstained on both controls. Grounded raw and application outputs selected checkout and abstained on both controls. Its fault explanation cites checkout's positive error metric and two ERROR spans. The selected pool also contains downstream payment ERROR spans, but the answer does not justify prioritizing checkout over payment. This is an injection-label non-match, not proven incorrect first action: no optimal-action annotation exists. The direct answer's claim that errors originate in payment also exceeds what nesting alone proves.

The normal grounded explanation again implies zero errors across all services, contradicted by selected load-generator records. Correct abstention therefore does not imply faithful explanation. A transparent full-pool rule reference abstained on development and both confirmation faults because checkout, frontend and payment tied at 13.5; it is not a matched 48-record comparator. The first

confirmation failed to repeat the development label match, so the candidate is not fully promoted. Triplet 004 subsequently failed in collection. A3 permitted one frozen replacement 005, with all failed and unattempted phases disclosed; it did not permit tuning.

Figure 5 appears at full width in Appendix C.

7.2 Replacement confirmation and final accounting

Replacement 005 completed its normal/fault/recovery sequence, using its own normal baseline. Its path checks respectively found 4/6/6 linked traces and 8/12/14 linked payment spans; all 12 fault payment spans were ERROR, versus zero linked payment ERROR spans in both controls. All six fixed calls completed. Grounded and direct raw/application decisions matched payment on the fault and deferred on both controls. Grounded cites the payment metric, payment child spans, and a checkout parent; its explanation supports an observed investigation lead without proving deepest cause. The direct explanation says the records confirm an initiating service, a stronger causal assertion than this observational analysis establishes.

Confirmation group	Direct: fault / controls	Grounded: fault / controls
003	payment; defer 2/2	checkout; defer 2/2
005	payment; defer 2/2	payment; defer 2/2

Table 6. Frozen confirmation raw and application decisions. Injection-label match is 2/2 for direct and 1/2 for grounded; both methods defer on 4/4 controls. These small, dependent counts do not score optimal first action or support same-task superiority. The 003 counterexample remains; 005 does not erase it.

The round stops at 30 calls, with 28 completed and two development HTTP 503 failures. No additional model tuning, quota escalation or retries are included. Both omitted-modality views remain untested. Capture failures, API availability, label matches, structural validation and semantic support are reported separately in the public aggregate.

Usage is known for 28 calls and unknown for two failures: 201,281 input tokens and 4,497 output tokens, 205,778 total. Latency is recorded for all 30 attempts (sum 219,724 ms, median 1,433.5 ms), including failed requests. Confirmation direct and grounded totals are 43,970 and 43,614 tokens, respectively. These are descriptive costs; billed currency, cache and thought-token usage remain unknown.

The 005 normal selected input has no ERROR spans or positive error-rate metrics, so its normality wording is not contradicted within that input. An unselected load-generator metric is positive, showing selection's visibility limit. Recovery 005 contains positive flagd/recommendation metrics despite all selected spans being OK. Without an anomaly threshold, those counters alone do not prove an active incident; the blanket absence-of-anomaly explanation is insufficiently justified. The actual parent citation on 005 fault is more complete than 002, but neither this improvement in the citation trail nor label match establishes reliable causal reasoning.

8. Discussion and next challenges

Human decisions supplied scope and scrutiny; AI supported schema design, code, debugging, and writing. Co-design corrected benchmark framing and clarified first inspection after the model requested deeper-cause evidence. The negative selector result and flawed control wording show why completed calls, label matches, valid citations, and useful investigation must remain distinct.

Limits include one injected mechanism, one host, dependent samples, evidence-conditioned acceptance, no logs, possible export delay, unresolved retention, and no on-call user study. Completed confirmation shows variable label agreement, not selector superiority, production readiness, or cross-company generalization. Challenge 3 should compare competing hypotheses using trace direction/timing and assess the next-check decision with human task evidence under a new frozen study. Challenge 4 should test predefined missing/delayed signals; Challenge 5 integrates verified features and observed review. Adding RAG or agents is not itself the objective.

9. Reproducibility, sources, and AI use

From the repository root, install `requirements.txt`, run `py -3.13 -m scripts.seed_official_examples`, and launch `py -3.13 -m streamlit run app.py --server.address 127.0.0.1 --server.port 8510`. Nine public official examples preserve exact incident/observation bytes and actual recorded responses, including the first confirmation's non-match. **Load recorded analysis** requires neither Docker nor an API key and makes no call. These are recorded observations, not a freshly running simulation. New official startup and capture require the dedicated WSL/Docker environment and `docs/SIMULATION_PROTOCOL.md`; run `py -3.13 -m scripts.check_shop` before capture. Preflight alone does not establish capture acceptance.

`scripts.run_paired_experiment` saves frozen inputs, responses, validators, and failures; `scripts.score_paired_experiment` separately reads private labels. The plan, A1/A2/A3 amendments, `selected_method.json`, claim audit, ledger, and runtime records specify provenance. Private configuration is excluded from inference/public releases. Before final example expansion, full local regression passed 186 tests in 34.27 seconds. The final nine-example public-source snapshot, without original cases or private API records, passed 185 tests in 39.17 seconds and skipped only the explicitly optional original-private-source audit; fresh seeding and verification passed. Its first attempt had 20 temporary-directory setup errors under sandbox permissions, with 165 tests passing. The unchanged suite passed after using a fresh writable workspace temporary directory. These checks establish software behavior, not diagnostic performance. The public repository contains the implementation, aggregate results, nine recorded examples and report. Published implementation commit `ec09fb3` was cloned into a fresh checkout on this configured Python environment: all 27 example JSON files, the PDF, and frozen schema/scorer byte hashes matched, and fresh seeding plus verification passed without Docker or API calls. Publication is separate from diagnostic or user-benefit validation.

Code is MIT; Shop remains Apache-2.0. Sources are the versioned papers in Section 3, pinned Shop, and the course handout. Codex assisted investigation, implementation, writing, and simulated review. `RESEARCH_REVIEW_20261001.md` and `CLAIM_AUDIT_20261001.md` record objections and unresolved limits; simulated agreement is not certification.

The bounded loop adapts Karpathy's autoresearch and a Codex-compatible skill, version 2.2.2. The original trains a small model; here the loop tests a fixed-budget inference pipeline. Reading and adapting that protocol does not imply installing its automation or training a model.

References

- [1] OpenRCA. ICLR 2025 conference paper, Section 2.3. [Paper](#) and [code/data entry point](#).
- [2] EviRCA. arXiv:2609.19825v1, September 17, 2026; preprint, Sections III-B/C and V-B. [Paper](#) and [implementation](#).
- [3] MicroRCA-Agent. arXiv:2509.15635v1, September 19, 2025; competition technical report, Sections 3.5 and 4.1. [Paper](#) and [implementation](#).
- [4] OpenTelemetry Demo 3.1.0. [Pinned source](#).
- [5] Andrej Karpathy. [autoresearch](#). Bounded iteration protocol adapted for this inference study.

Appendix A. Course requirement coverage

Handout item	Coverage in this report	Remaining evidence requirement
1. Problem / P1	Sections 1-2; Figure 1	User benefit unmeasured
2. Human Design	Section 4.1; Figure 1	Honest reconstruction disclosed
3. AI Design	Section 4.2; Figure 2	Proposal summary; no invented transcript
4. Co-design	Section 4.3; Table 1	Changes traceable to records
5. Project data	Section 5; Table 2	Nine modeled phase windows; one mechanism
6. Methods / architecture	Sections 5-6; Figure 3	A1/A2 development and A3 replacement disclosed
7. Grounding / provenance	Section 7; Table 4	Simulated audit; semantic failures retained
8. End-to-end feature	Section 5; Figure 4	Actual official evidence UI; user benefit unmeasured
9. Results / 5-10 cases	Section 7; Tables 2-4 and 6; supplement below	Nine dependent official phase windows, not diverse independent faults
10. Three-stage comparison	Table 1, all 11 aspects	User-quality comparison unmeasured
11. GitHub / reproducibility	Section 9	Public commit and fresh-checkout replay verified
12. Lessons / next challenge	Section 8	Fresh controlled evidence required

Supplementary historical native development

These five earlier four-process Python cases used 24 requests per phase, concurrency six, and 534 observations. They are historical development evidence, not official-Shop data or confirmation.

Native case	Expected service	Strong direct	Grounded chronological	Grounded prioritized
Payment capacity	payment	payment	payment	payment
Catalog error	catalog	catalog	catalog	catalog
Checkout error	checkout	checkout	checkout	checkout
Clean control	none	abstain	abstain	abstain
Payment delay	payment	checkout	checkout	abstain

Table 5. Historical native outcomes. Model conditions each localized 3/4 faults; the rule baseline localized 4/4. All abstained on the clean case. The post-hoc selection/order ablation avoided one wrong attribution by abstaining, without localization gain. It is not independent confirmation or official evidence.

Appendix B. Full-width comparison tables

Aspect	Human Design	AI Design	Reviewed co-design
Problem understanding	First investigation priority	Broad incident diagnosis	Checkout symptom propagation
Data understanding	Public incident records	Benchmark adapter	Captured Shop; private interventions
Architecture	One complete P1	Generic analysis pipeline	Bounded evidence-to-review path
Multimodal design	Metrics, traces, available logs	Compact representations	Preserve values, times, trace links
Retrieval / RAG	Relevant records	Evidence selection	Balanced rule/metadata policy; no embeddings
Models / tools	Reuse existing capability	Python plus LLM	Tools compute; Gemini interprets
Grounding	Human source inspection	Evidence IDs	Resolve IDs; separately audit support
Implementation	Local review interface	Code/schema assistance	Streamlit plus answer-isolated contract
Debugging	Verify assumptions	Suggested repairs	Record runtime/parser fixes separately
Evaluation	Actual cases and failures	Initial prompt comparison	Same inputs; raw/validated outputs; frozen confirmation
Final quality	Understandable decision aid	Runnable proposal	Inspectable prototype; benefit unmeasured

Table 1. Eleven-aspect course comparison. Design improvements are arguments about traceability and scope, not measured performance gains.

Condition / prompt	Completed / scheduled	Fault payment agreement: raw -> app	Core controls abstaining / completed	Unavailable
Prioritized / direct	3/3	1/1 -> 1/1	2/2	None
Prioritized / grounded	2/3	1/1 -> 0/1	1/1	Normal 503
Balanced / direct	3/3	1/1 -> 1/1	2/2	None
Balanced / grounded	3/3	0/1 -> 0/1	2/2	None
Priority / direct	2/3	1/1 -> 1/1	1/1	Recovery 503
Priority / grounded	3/3	1/1 -> 1/1	2/2	None

Table 3. Official development decisions. "Priority" changes the grounded task to first inspection with balanced inputs. Agreement uses evaluator-only injection labels; it is not best-action correctness. Missing controls remain unknown. Original grounded fault names payment raw but validation defers because citations are spans only; balanced grounded already abstains raw. Priority grounded cites candidate-local metric/spans and names payment.

Appendix C. Recorded application and runtime views




Figure 4. Actual recorded analysis of an accepted official fault case. The interface displays the saved payment inspection priority, explanation, and cited IDs, explicitly identifying replay and distinguishing priority from proven root cause. The screenshot establishes interface behavior, not diagnostic correctness or user benefit.




Figure 5. Actual first-confirmation non-match, replayed without a new API call. The model prioritizes checkout despite the payment intervention. Its cited checkout errors are real, but the explanation does not establish priority over the downstream payment alternative. No optimal-action annotation is available.



The best telescopes to see the world closer

Go Shopping

Hot Products



Solar System Color Imager
\$ 175.00



Recorded official Shop storefront on October 1, 2026. This is runtime visual evidence, not a diagnostic score.